

Utilize your CPU power

Cache optimizations and SIMD instructions

Mario Mulansky

ISC-CNR, Institute for Complex Systems, Florence, Italy




Utilize your CPU power

Cache optimizations and SIMD instructions

Mario Mulansky

ISC-CNR, Institute for Complex Systems, Florence, Italy



Performance

Numerical Simulations:

- Flops Flops Flops
- Actually: Simulation runtime
- Not: DB queries / Requests per second, Response rate, ...

Performance

Numerical Simulations:

- Flops Flops Flops
- Actually: Simulation runtime
- Not: DB queries / Requests per second, Response rate, ...

Performance bounds:

- CPU power – Floating point operations per second: Flops/s
- Memory bandwidth

Performance

Numerical Simulations:

- Flops Flops Flops
- Actually: Simulation runtime
- Not: DB queries / Requests per second, Response rate, ...

Performance bounds:

- CPU power – Floating point operations per second: Flops/s
- Memory bandwidth

This talk:

- Bandwidth bottleneck
- Example: Numerical algorithm (ODE solver)
- Bandwidth bound \rightarrow Flops bound
- More speed: Boost.SIMD

Data bandwidth and latency limitation

Modern CPUs: ~ 3 GHz, 1 Op/cycle $\rightarrow 3$ GFlops/s

Data bandwidth and latency limitation

Modern CPUs: ~ 3 GHz, 1 Op/cycle $\rightarrow 3$ GFlops/s

Problem: Data transfer from memory to the CPU \sim : 16 GB/s

Example $\mathbf{x} = \mathbf{a} + \mathbf{b}$: 24 Bytes/Operation

3 GFlops/s \leftrightarrow 72 GByte/s or 16 GByte/s \leftrightarrow 0.7 GFlop/s

Data bandwidth and latency limitation

Modern CPUs: ~ 3 GHz, 1 Op/cycle $\rightarrow 3$ GFlops/s

Problem: Data transfer from memory to the CPU \sim : 16 GB/s

Example $\mathbf{x} = \mathbf{a} + \mathbf{b}$: 24 Bytes/Operation


3 GFlops/s \leftrightarrow 72 GByte/s or 16 GByte/s \leftrightarrow 0.7 GFlop/s

Cache

Intel Sandy Bridge

Cache	Size	GB/s	Cycles
L1	32K	350	4
L2	256K	250	12
L3	3-6M	100	12-30
RAM	4-32G	16	20-100

Intel® Core™ i7-3960X Processor Die Detail



Cache effect: Ordinary Differential Equations

Solve a system of N small, independent ODEs: $\dot{r}_i = f(r_i, t)$

(Parameter study, Monte-Carlo)

Cache effect: Ordinary Differential Equations

Solve a system of N small, independent ODEs: $\dot{r}_i = f(r_i, t)$

(Parameter study, Monte-Carlo)

Two approaches:

- time first: $r_1(t) \rightarrow r_1(t + \Delta t) \rightarrow r_1(t + 2\Delta t) \dots$, then $r_2 \dots$
- vector first: $\mathbf{r}(t) \rightarrow \mathbf{r}(t + \Delta t) \rightarrow \mathbf{r}(t + 2\Delta t) \dots$


Cache effect: Ordinary Differential Equations

Solve a system of N small, independent ODEs: $\dot{r}_i = f(r_i, t)$

(Parameter study, Monte-Carlo)

Two approaches:

- time first: $r_1(t) \rightarrow r_1(t + \Delta t) \rightarrow r_1(t + 2\Delta t) \dots$, then $r_2 \dots$
- vector first: $\mathbf{r}(t) \rightarrow \mathbf{r}(t + \Delta t) \rightarrow \mathbf{r}(t + 2\Delta t) \dots$




Cache effect: Ordinary Differential Equations

Solve a system of N small, independent ODEs: $\dot{r}_i = f(r_i, t)$

(Parameter study, Monte-Carlo)

Two approaches:

- time first: $r_1(t) \rightarrow r_1(t + \Delta t) \rightarrow r_1(t + 2\Delta t) \dots$, then $r_2 \dots$
- vector first: $\mathbf{r}(t) \rightarrow \mathbf{r}(t + \Delta t) \rightarrow \mathbf{r}(t + 2\Delta t) \dots$



ODE with nearest neighbor coupling

$$\begin{aligned}\dot{r}_i &= f_i(\mathbf{r}, t) \\ &= \underbrace{h_i(r_i, t)}_{\text{local}} + \underbrace{g_i(r_i, r_{i-1}, r_{i+1}, t)}_{\text{n. n. coupling}}\end{aligned}$$

ODE with nearest neighbor coupling

$$\begin{aligned}\dot{r}_i &= f_i(\mathbf{r}, t) \\ &= \underbrace{h_i(r_i, t)}_{\text{local}} + \underbrace{g_i(r_i, r_{i-1}, r_{i+1}, t)}_{\text{n. n. coupling}}\end{aligned}$$

Time discretization \mathbf{r}_t with Δt .

R-K algorithms: $\mathbf{r}_t \rightarrow \mathbf{r}_{t+\Delta t}$
with s stages $j = 1 \dots s$.

at each stage:

$$\begin{aligned}\mathbf{k}^j &= \mathbf{f}(\mathbf{r}', t') \\ \mathbf{r}' &= \mathbf{r}_t + \sum_{n < j} a_{j,n} \mathbf{k}^n \Delta t\end{aligned}$$

ODE with nearest neighbor coupling


$$\begin{aligned}\dot{r}_i &= f_i(\mathbf{r}, t) \\ &= \underbrace{h_i(r_i, t)}_{\text{local}} + \underbrace{g_i(r_i, r_{i-1}, r_{i+1}, t)}_{\text{n. n. coupling}}\end{aligned}$$

Time discretization \mathbf{r}_t with Δt .

R-K algorithms: $\mathbf{r}_t \rightarrow \mathbf{r}_{t+\Delta t}$
with s stages $j = 1 \dots s$.

at each stage:

$$\begin{aligned}\mathbf{k}^j &= \mathbf{f}(\mathbf{r}', t') \\ \mathbf{r}' &= \mathbf{r}_t + \sum_{n < j} a_{j,n} \mathbf{k}^n \Delta t\end{aligned}$$



ODE with nearest neighbor coupling

$$\begin{aligned}\dot{r}_i &= f_i(\mathbf{r}, t) \\ &= \underbrace{h_i(r_i, t)}_{\text{local}} + \underbrace{g_i(r_i, r_{i-1}, r_{i+1}, t)}_{\text{n. n. coupling}}\end{aligned}$$


Time discretization \mathbf{r}_t with Δt .

R-K algorithms: $\mathbf{r}_t \rightarrow \mathbf{r}_{t+\Delta t}$
with s stages $j = 1 \dots s$.

at each stage:


$$\begin{aligned}\mathbf{k}^j &= \mathbf{f}(\mathbf{r}', t') \\ \mathbf{r}' &= \mathbf{r}_t + \sum_{n < j} a_{j,n} \mathbf{k}^n \Delta t\end{aligned}$$

Problem: Coupling prevents
“time first” approach




Improve Cache Usage: Granularity

Coupling: only “vector first”?



Improve Cache Usage: Granularity


Coupling: only “vector first”?




Clustering!

Improve Cache Usage: Granularity

Coupling: only “vector first”?




Clustering!



Improve Cache Usage: Granularity


Coupling: only “vector first”?



Clustering!


s iterations for each cluster at once \rightarrow better cache usage

Price: additional overlap computations



Improve Cache Usage: Granularity

Coupling: only “vector first”?




Clustering!

s iterations for each cluster at once \rightarrow better cache usage


Price: additional overlap computations

Optimal granularity?



Improve Cache Usage: Granularity

Coupling: only “vector first”?




Clustering!

s iterations for each cluster at once \rightarrow better cache usage

Price: additional overlap computations

Optimal granularity?

\sim Cache size... Measure!



Performance Results

Coupled Rössler systems, $N = 2^{20} \approx 10^6$ (24 MB)

Intel Xeon E5-2690 @ 3.8GHz, Intel Compiler 15.0.0


Performance measurements with `likwid` framework

Performance Results

Coupled Rössler systems, $N = 2^{20} \approx 10^6$ (24 MB)

Intel Xeon E5-2690 @ 3.8GHz, Intel Compiler 15.0.0

Performance measurements with `likwid` framework





Performance Results

Coupled Rössler systems, $N = 2^{20} \approx 10^6$ (24 MB)

Intel Xeon E5-2690 @ 3.8GHz, Intel Compiler 15.0.0

Performance measurements with `likwid` framework





Performance Results

Coupled Rössler systems, $N = 2^{20} \approx 10^6$ (24 MB)

Intel Xeon E5-2690 @ 3.8GHz, Intel Compiler 15.0.0


Performance measurements with `likwid` framework



Bandwidth bound \rightarrow Flops/s bound

Increase Flops/s: SIMD instructions


SIMD: Single Instruction Multiple Data



Additional registers and instructions in modern CPUs

Increase Flops/s: SIMD instructions

SIMD: Single Instruction Multiple Data



Additional registers and instructions in modern CPUs

Compilers try to use those automatically

Often, explicit SIMD code improve Flops/s significantly

Only helpful, if algorithm is Flops bound!

- Abstraction of SIMD instructions:
(SSE3, SSE4.1, SSE4.2, AVX, FMA4, AltiVec, Intel MIC).
- Fundamental ingredient: SIMD registers \rightarrow SIMD pack.
- Expression template for optimization possibilities.
- Keep container–iterator–algorithm abstraction.

- Abstraction of SIMD instructions:
(SSE3, SSE4.1, SSE4.2, AVX, FMA4, AltiVec, Intel MIC).
- Fundamental ingredient: SIMD registers \rightarrow SIMD pack.
- Expression template for optimization possibilities.
- Keep container–iterator–algorithm abstraction.

`simd::pack<T>`

- `pack<T,N>` SIMD register with `N` elements of `T`.
- `pack<T>` **automatically chooses `N` from available hardware.**
- `T` must be integral, e.g. `int`, `float`, `double`...
- `N` must be power of 2.

- Abstraction of SIMD instructions:
(SSE3, SSE4.1, SSE4.2, AVX, FMA4, AltiVec, Intel MIC).
- Fundamental ingredient: SIMD registers \rightarrow SIMD pack.
- Expression template for optimization possibilities.
- Keep container–iterator–algorithm abstraction.

`simd::pack<T>`

- `pack<T,N>` SIMD register with `N` elements of `T`.
- `pack<T>` **automatically chooses `N` from available hardware.**
- `T` must be integral, e.g. `int`, `float`, `double`...
- `N` must be power of 2.
- Operations available `+`, `-`, `*`, `/`
- Math functions available: `pow`, `abs`, `sqrt`, `log10`, ...

Quick Example: vector addition $\mathbf{x} = \alpha \cdot \mathbf{a} + \mathbf{b}$

```
typedef vector<double> vec;  
  
double alpha;  
vec x(N), a(N), b(N);  
  
for(int n=0; n<x.size(); ++n)  
    x[n] = alpha * a[n] + b[n];
```


Quick Example: vector addition $\mathbf{x} = \alpha \cdot \mathbf{a} + \mathbf{b}$

```
typedef vector<double> vec;  
  
double alpha;  
vec x(N), a(N), b(N);  
  
for(int n=0; n<x.size(); ++n)  
    x[n] = alpha * a[n] + b[n];
```

```
typedef simd::pack<double> pack; // automatic size  
typedef vector<pack, simd::allocator<pack> > vec;  
static const size_t pack_size = pack::static_size;  
static const int M = N/pack_size;  
  
double alpha;  
vec x(M), a(M), b(M);  
  
for(int n=0; n<x.size(); ++n)  
    x[n] = alpha * a[n] + b[n];
```

Algorithm does not change

```
for(int n=0; n<x.size(); ++n)  
    x[n] = alpha * a[n] + b[n];
```

even more clear:

```
transform(a, b, x, [alpha](double a_n, double b_n)  
    {return alpha * a_n + b_n;})
```

Algorithm does not change

```
for(int n=0; n<x.size(); ++n)  
    x[n] = alpha * a[n] + b[n];
```

even more clear:

```
transform(a, b, x, [alpha](double a_n, double b_n)  
    {return alpha * a_n + b_n;})
```

With Boost.SIMD: change typedefs, don't touch algorithm.

SIMD for ODE Simulation

ODE iteration $\mathbf{r}_t \rightarrow \mathbf{r}_{t+\Delta t}$: some sort of transform

Abstraction Boost.odeint provides:

- generic algorithms
- container independent implementation
- exchangeable backends

```
typedef vector<double> state_type;  
  
state_type x(N);  
odeint::runge_kutta4<state_type> rk4;  
odeint::integrate_const(rk4, roessler, x, 0.0, T, dt);
```

SIMD for ODE Simulation

ODE iteration $\mathbf{r}_t \rightarrow \mathbf{r}_{t+\Delta t}$: some sort of transform

Abstraction Boost.odeint provides:

- generic algorithms
- container independent implementation
- exchangeable backends


```
typedef simd::pack<double> simd_pack;  
typedef vector<simd_pack,  
             simd::allocator<simd_pack> > state_type;  
static const size_t pack_size = simd_pack::static_size;  
static const int M = N/pack_size;  
  
state_type x(M);  
odeint::runge_kutta4<state_type> rk4;  
odeint::integrate_const(rk4, roessler, x, 0.0, T, dt);
```

SIMD Performance Results

Boost.odeint + Boost.SIMD



SIMD Performance Results

Boost.odeint + Boost.SIMD




SIMD Performance Results

Boost.odeint + Boost.SIMD



Size Dependence



Summary + Conclusions

- Granularity \rightarrow data transfer \searrow , BW bound \rightarrow Flops/s bound.
- Increase Op/cycle via SIMD \rightarrow total performance gain 3x.
- Known for stencil computations: space- and time-blocking.

Summary + Conclusions

- Granularity \rightarrow data transfer \searrow , BW bound \rightarrow Flops/s bound.
- Increase Op/cycle via SIMD \rightarrow total performance gain 3x.
- Known for stencil computations: space- and time-blocking.

Take-home-message:

- Data size $>$ L2 cache size \rightarrow **introduce granularity.**
- Write generic algorithms, substitute Boost.SIMD.
- Source code: <https://github.com/mariomulansky/olsos>
- C++ with Boost.odeint and Boost.SIMD \sim 200 lines of code.

Summary + Conclusions

- Granularity \rightarrow data transfer \searrow , BW bound \rightarrow Flops/s bound.
- Increase Op/cycle via SIMD \rightarrow total performance gain 3x.
- Known for stencil computations: space- and time-blocking.

Take-home-message:

- Data size $>$ L2 cache size \rightarrow **introduce granularity.**
- Write generic algorithms, substitute Boost.SIMD.
- Source code: <https://github.com/mariomulansky/olsos>
- C++ with Boost.odeint and Boost.SIMD \sim 200 lines of code.



www.odeint.com